

## 4.1.2 LSTM Time Series to Predict Ground Level PM10 with Lagged Variables

### Chapter 4.1.2

```
1 from google.colab import drive
2 drive.mount('/content/gdrive')
```

Mounted at /content/gdrive

```

1 # Import librries
2 import pandas as pd
3 from pandas import DataFrame
4 from pandas import concat
5 from sklearn.preprocessing import LabelEncoder
6 from sklearn.preprocessing import MinMaxScaler
7 import glob #for maps
8 import plotly.graph_objects as go
9 import plotly.offline
10 import matplotlib.pyplot as plt
11 from datetime import datetime
12 import warnings
13 import itertools
14 import numpy as np
15 from statsmodels.tsa.stattools import adfuller
16 from sklearn.metrics import mean_squared_error, r2_score, mean_absolute_error
17 from math import sqrt
18 warnings.filterwarnings("ignore")
19 import plotly.express as px
20 import seaborn as sns
21 import pandas as pd
22 import plotly.graph_objects as go
23 import matplotlib.pyplot as plt
24 %matplotlib inline
25
26
27 from sklearn.model_selection import train_test_split
28 from tensorflow.keras.layers import Dense, Dropout, LSTM
29 import keras.models
30 import tensorflow
31 from keras.models import Sequential
32 from tensorflow.keras.layers import LSTM
33 from tensorflow.keras.layers import Dense

```

## ▼ Data Reading and Additional Processing

```
1 fin_data = pd.read_csv('/content/gdrive/MyDrive/Sent/final_data_May20_Sept22.
```

```
1 ar', 'Month', 'Day', 'Hour']][ 'nitrogendioxide_tropospheric_column', 'nitrogend
```

```
1 fin_data_exp=fin_data_exp.reset_index()
2 # fin_data_exp=fin_data_exp.drop(columns=['level_0', 'index', 'Year', 'Month'])
3 # fin_data_exp.head()
```

```
1 fin_data_exp=fin_data_exp.drop(columns=['Year', 'Month', 'Day', 'Hour'])
2 fin_data_exp.head()
```

	nitrogendioxide_tropospheric_column	nitrogendioxide_tropospheric_column
0	0.000008	
1	0.000017	
2	0.000046	
3	0.000033	
4	0.000039	

## ▼ Feature Engineering for LSTM

### Reference:

The model scrip was adopted from the following webpage:

<https://machinelearningmastery.com/multivariate-time-series-forecasting-lstms-keras/>

<https://machinelearningmastery.com/time-series-prediction-lstm-recurrent-neural-networks-python-keras/>

<https://datascience.stackexchange.com/questions/76826/why-are-predictions-from-my-lstm-neural-network-lagging-behind-true-values>

```
1
2 # convert series to supervised learning
3 def series_to_supervised(data, n_in=1, n_out=1, dropnan=True):
4     n_vars = 1 if type(data) is list else data.shape[1]
5     df = DataFrame(data)
6     cols, names = list(), list()
7     # input sequence (t-n, ... t-1)
8     for i in range(n_in, 0, -1):
9         cols.append(df.shift(i))
```

```
10     names += [('var%d(t-%d)' % (j+1, i)) for j in range(n_vars)]
11     # forecast sequence (t, t+1, ... t+n)
12     for i in range(0, n_out):
13         cols.append(df.shift(-i))
14         if i == 0:
15             names += [('var%d(t)' % (j+1)) for j in range(n_vars)]
16         else:
17
18             names += [('var%d(t+%d)' % (j+1, i)) for j in range(n_vars)]
19     # put it all together
20     agg = concat(cols, axis=1)
21     agg.columns = names
22     # drop rows with NaN values
23     if dropnan:
24         agg.dropna(inplace=True)
25     return agg
26
27 # load dataset
28
29 values = fin_data_exp.values
30 # integer encode direction
31 encoder = LabelEncoder()
32
33
34 # ensure all data is float
35 values = values.astype('float32')
36 # normalize features
37
38 #values = lstm_data(values, 4)
39
40 scaler = MinMaxScaler(feature_range=(0, 1))
41 scaled = scaler.fit_transform(values)
42
43
44 # frame as supervised learning
45 reframed = series_to_supervised(scaled, 1, 1)
46 #reframed.drop(reframed.columns[[9,10,11]], axis=1, inplace=True)
47 #reframed.drop(reframed.columns[[6,7,8,9,10]], axis=1, inplace=True)
48 #reframed.drop(reframed.columns[[6,7,8,9,10]], axis=1, inplace=True) #window
49 reframed.head()
50
```

	<code>var1(t-1)</code>	<code>var2(t-1)</code>	<code>var3(t-1)</code>	<code>var4(t-1)</code>	<code>var5(t-1)</code>	<code>var6(t-1)</code>	<code>var1(t)</code>	<code>var2(t)</code>	<code>var3(t)</code>
1	0.155634	0.013603	0.015260	0.427553	0.546352	0.129479	0.164217	0.005729	0.0
2	0.164217	0.005729	0.005855	0.654652	0.739136	0.129479	0.191809	0.036969	0.0
3	0.191809	0.036969	0.033096	0.426778	0.564740	0.201648	0.179365	0.018086	0.0
4	0.179365	0.018086	0.015359	0.500033	0.553890	0.201648	0.185115	0.034446	0.0
5	0.185115	0.034446	0.032485	0.407516	0.745865	0.373201	0.188677	0.039848	0.0

```
1 reframed.shape
```

```
(985, 18)
```

```
1 reframed1 = reframed.iloc[:, :-1]
2 reframed1 = lstm_data(reframed1, 3)
```

```
1 reframed.shape
```

```
(985, 18)
```

```
1 values1.shape
```

```
(985, 18)
```

## ▼ **Model 1 LSTM (window one)**

Double-click (or enter) to edit

```

1 values = reframed.values
2 #values1 = reframed.values
3 x, x_test, y, y_test = train_test_split(values[:, :-1], values[:, -1], test_size=
4 x_train, x_val, train_y, val_y = train_test_split(x, y, test_size = 0.2, train_s
5 # reshape input to be 3D [samples, timesteps, features]
6
7 train_X = x_train.reshape((x_train.shape[0], 1, x_train.shape[1]))
8 val_X = x_val.reshape((x_val.shape[0], 1, x_val.shape[1]))
9 test_X = x_test.reshape((x_test.shape[0], 1, x_test.shape[1]))
10 print(train_X.shape, train_y.shape, val_X.shape, val_y.shape, test_X.shape, y

(709, 1, 6) (709,) (178, 1, 6) (178,) (99, 1, 6) (99,)

```

```

1
2 model = Sequential()
3 model.add(LSTM(500, input_shape=(train_X.shape[1], train_X.shape[2])))
4 #model.add(LSTM(1000, input_shape=(train_X.shape[1], train_X.shape[2])))
5 #model.add(LSTM(2000, input_shape=(train_X.shape[1], train_X.shape[2])))
6 model.add(Dropout(0.2))
7 model.add(Dense(1))
8
9 model.compile(loss='mae', optimizer='adam') #mean absolute error
10
11
12
13 history = model.fit(train_X, train_y, epochs=50, batch_size=64, validation_da
14 # plot history
15 plt.figure(figsize=(15, 10))
16
17 pyplot.plot(history.history['loss'], label='train')
18 pyplot.plot(history.history['val_loss'], label='val')
19 pyplot.legend()
20 pyplot.show()
21
22 y_hat = model.predict(train_X)
23 y_hat_val = model.predict(val_X)
24
25 plt.figure(figsize=(15, 10))
26 pyplot.plot(y_hat, label='predict', color = 'red')
27 pyplot.plot(train_y, label='true', color = 'yellow')
28 pyplot.legend()
29
30 pyplot.show()
31
32
33 plt.figure(figsize=(15, 10))
34 pyplot.plot(y_hat_val, label='predict', color = 'blue')

```

```

35 pyplot.plot(val_y, label='true', color = 'pink')
36 pyplot.legend()
37
38 pyplot.show()
39
40

```


```

Epoch 1/50
12/12 - 3s - loss: 0.4691 - val_loss: 0.3270 - 3s/epoch - 219ms/step
Epoch 2/50
12/12 - 0s - loss: 0.2375 - val_loss: 0.1824 - 246ms/epoch - 20ms/step
Epoch 3/50
12/12 - 0s - loss: 0.1870 - val_loss: 0.1870 - 243ms/epoch - 20ms/step
Epoch 4/50
12/12 - 0s - loss: 0.1936 - val_loss: 0.1725 - 259ms/epoch - 22ms/step
Epoch 5/50
12/12 - 0s - loss: 0.1771 - val_loss: 0.1688 - 239ms/epoch - 20ms/step
Epoch 6/50
12/12 - 0s - loss: 0.1712 - val_loss: 0.1649 - 246ms/epoch - 20ms/step
Epoch 7/50
12/12 - 0s - loss: 0.1695 - val_loss: 0.1612 - 256ms/epoch - 21ms/step
Epoch 8/50
12/12 - 0s - loss: 0.1641 - val_loss: 0.1562 - 255ms/epoch - 21ms/step
Epoch 9/50
12/12 - 0s - loss: 0.1581 - val_loss: 0.1527 - 245ms/epoch - 20ms/step
Epoch 10/50
12/12 - 0s - loss: 0.1550 - val_loss: 0.1484 - 243ms/epoch - 20ms/step
Epoch 11/50
12/12 - 0s - loss: 0.1515 - val_loss: 0.1441 - 241ms/epoch - 20ms/step
Epoch 12/50
12/12 - 0s - loss: 0.1471 - val_loss: 0.1394 - 282ms/epoch - 23ms/step
Epoch 13/50
12/12 - 0s - loss: 0.1427 - val_loss: 0.1350 - 242ms/epoch - 20ms/step
Epoch 14/50
12/12 - 0s - loss: 0.1394 - val_loss: 0.1313 - 260ms/epoch - 22ms/step
Epoch 15/50
12/12 - 0s - loss: 0.1372 - val_loss: 0.1291 - 240ms/epoch - 20ms/step
Epoch 16/50
12/12 - 0s - loss: 0.1346 - val_loss: 0.1280 - 250ms/epoch - 21ms/step
Epoch 17/50
12/12 - 0s - loss: 0.1313 - val_loss: 0.1278 - 244ms/epoch - 20ms/step
Epoch 18/50
12/12 - 0s - loss: 0.1305 - val_loss: 0.1278 - 246ms/epoch - 20ms/step
Epoch 19/50
12/12 - 0s - loss: 0.1286 - val_loss: 0.1273 - 258ms/epoch - 21ms/step
Epoch 20/50
12/12 - 0s - loss: 0.1291 - val_loss: 0.1273 - 253ms/epoch - 21ms/step
Epoch 21/50
12/12 - 0s - loss: 0.1275 - val_loss: 0.1276 - 233ms/epoch - 19ms/step
Epoch 22/50
12/12 - 0s - loss: 0.1305 - val_loss: 0.1276 - 244ms/epoch - 20ms/step
Epoch 23/50
12/12 - 0s - loss: 0.1294 - val_loss: 0.1283 - 247ms/epoch - 21ms/step

```


```
Epoch 24/50
12/12 - 0s - loss: 0.1299 - val_loss: 0.1292 - 252ms/epoch - 21ms/step
Epoch 25/50
12/12 - 0s - loss: 0.1277 - val_loss: 0.1278 - 261ms/epoch - 22ms/step
Epoch 26/50
12/12 - 0s - loss: 0.1287 - val_loss: 0.1280 - 244ms/epoch - 20ms/step
Epoch 27/50
12/12 - 0s - loss: 0.1285 - val_loss: 0.1278 - 252ms/epoch - 21ms/step
Epoch 28/50
12/12 - 0s - loss: 0.1283 - val_loss: 0.1274 - 254ms/epoch - 21ms/step
Epoch 29/50
12/12 - 0s - loss: 0.1290 - val_loss: 0.1283 - 234ms/epoch - 20ms/step
Epoch 30/50
12/12 - 0s - loss: 0.1280 - val_loss: 0.1285 - 244ms/epoch - 20ms/step
Epoch 31/50
12/12 - 0s - loss: 0.1281 - val_loss: 0.1283 - 242ms/epoch - 20ms/step
Epoch 32/50
12/12 - 0s - loss: 0.1267 - val_loss: 0.1273 - 247ms/epoch - 21ms/step
Epoch 33/50
12/12 - 0s - loss: 0.1269 - val_loss: 0.1273 - 247ms/epoch - 21ms/step
Epoch 34/50
12/12 - 0s - loss: 0.1276 - val_loss: 0.1271 - 240ms/epoch - 20ms/step
Epoch 35/50
12/12 - 0s - loss: 0.1276 - val_loss: 0.1290 - 239ms/epoch - 20ms/step
Epoch 36/50
12/12 - 0s - loss: 0.1285 - val_loss: 0.1280 - 251ms/epoch - 21ms/step
Epoch 37/50
12/12 - 0s - loss: 0.1278 - val_loss: 0.1287 - 244ms/epoch - 20ms/step
Epoch 38/50
12/12 - 0s - loss: 0.1278 - val_loss: 0.1271 - 246ms/epoch - 20ms/step
Epoch 39/50
12/12 - 0s - loss: 0.1289 - val_loss: 0.1283 - 247ms/epoch - 21ms/step
Epoch 40/50
12/12 - 0s - loss: 0.1284 - val_loss: 0.1269 - 258ms/epoch - 22ms/step
Epoch 41/50
12/12 - 0s - loss: 0.1275 - val_loss: 0.1289 - 247ms/epoch - 21ms/step
Epoch 42/50
12/12 - 0s - loss: 0.1279 - val_loss: 0.1270 - 245ms/epoch - 20ms/step
Epoch 43/50
12/12 - 0s - loss: 0.1276 - val_loss: 0.1271 - 254ms/epoch - 21ms/step
Epoch 44/50
12/12 - 0s - loss: 0.1274 - val_loss: 0.1268 - 250ms/epoch - 21ms/step
Epoch 45/50
12/12 - 0s - loss: 0.1268 - val_loss: 0.1270 - 253ms/epoch - 21ms/step
Epoch 46/50
12/12 - 0s - loss: 0.1287 - val_loss: 0.1279 - 244ms/epoch - 20ms/step
Epoch 47/50
12/12 - 0s - loss: 0.1275 - val_loss: 0.1268 - 250ms/epoch - 21ms/step
Epoch 48/50
12/12 - 0s - loss: 0.1290 - val_loss: 0.1274 - 264ms/epoch - 22ms/step
Epoch 49/50
12/12 - 0s - loss: 0.1295 - val_loss: 0.1281 - 244ms/epoch - 20ms/step
Epoch 50/50
12/12 - 0s - loss: 0.1278 - val_loss: 0.1298 - 246ms/epoch - 20ms/step
```






23/23 [=====] - 1s 4ms/step


6/6 [=====] - 0s 5ms/step





```
1 #Test set
2
3 y_hat_test = model.predict(test_X)
4
5 plt.figure(figsize=(15, 10))
6
7 pyplot.plot(y_hat_test[75:90], label='predict', color = 'blue')
8 pyplot.plot(y_test, label='true', color = 'pink')
9 pyplot.legend()
10
11 pyplot.show()
12
13
14
15 print('MSE comparison ----')
16 print(f"LSTM:\t{mean_squared_error(y_hat_test, y_test):.4f}")
17 print('R2_score ----')
18 print(f"LSTMf:\t{r2_score(y_hat_test, y_test):.4f}")
19
```

4/4 [=====] - 0s 5ms/step



MSE comparison ----

LSTM: 0.0320

R2\_score ----

LSTMf: 0.1574